

DataChat

Agentic natural-language analytics over public open data.
Architecture, evaluation, and security.

Live demo: data-chat-seven.vercel.app

Source: github.com/sahil7359/DataChat

Generated 2026-08-14 from the repository documentation.

Contents

- 1. Overview** [README.md](#)
- 2. Architecture and Flow** [FLOW.md](#)
- 3. Sequence Diagrams** [AppFlow.md](#)
- 4. Data Model** [Schema.md](#)
- 5. Security** [SECURITY.md](#)
- 6. Change Log and Reasoning** [LEARN.md](#)

1. Overview

What the system does, the deliberately narrow scope it runs on, and the measured numbers it publishes. *Source: README.md*

<div align="center">

Ask open data in plain English. Get safe, verified SQL, a grounded answer, and a chart. An agentic natural-language analytics platform over public datasets - built as a real LangGraph agent, evaluated, observable, and secured, running entirely on free tiers.

!Python !FastAPI !LangGraph !MLflow !Next.js !Postgres !Cost !License

> [Open the live demo](#)

[**Hero GIF goes here**] - 15-20s: question -> plan -> SQL -> table -> chart. Instructions in docs/README.md (docs/README.md).

</div>

What it covers, and what it doesn't. 15 countries - GDP per capita, population and life expectancy (World Bank, 2022) - CO₂ emissions (Our World in Data, 2021-2022). It cannot answer about any other country, indicator or year, and it says so rather than guessing. **The narrowness is a cost decision, not an oversight** - the whole system runs at **\$0/month** on free tiers, and the point of the project is retrieval quality and evaluation, not corpus size. Widening it is a config change, not a rewrite.

Measured on Groq llama-3.3-70b-versatile, temperature=0	Score	n
Execution accuracy - result set equals the gold query's, exactly	0.810	21
Refusal accuracy - declined instead of inventing an answer	1.00	5
SQL valid rate - parses and passes the AST guardrail	0.952	21
Explanation faithfulness - prose grounded in the returned rows	0.905	21

Reproduce with `make eval-real`. Baselines are committed per provider in `backend/eval_baseline.json` (`backend/eval_baseline.json`); see Evaluation for how the gate works and Known limitations for what these numbers do *not* cover.

First load takes ~50s if the demo has been idle - the free backend sleeps after 15 min and the database scales to zero. A keep-warm ping every 12 minutes mitigates it.

Or skip the UI entirely:

```
curl -N -X POST https://datachat-api-wmpd.onrender.com/api/v1/chat \  
-H "Content-Type: application/json" \  
-d '{"question":"Which 5 countries had the highest CO2 per capita in 2022?"}'
```

[API](#) - [API docs](#) - [Architecture](#) - [Flow](#) (./FLOW.md) - [Design](#) (./Design.md) - [Changelog + reasoning](#) (./LEARN.md)

Technical overview (PDF, 30pp) (docs/DataChat-Technical-Overview.pdf) - architecture, flow, data model, security and change log in one document. Regenerate with `python docs/build_pdf.py`.

The problem & why it matters

Public open datasets are rich but locked behind SQL and BI tools. Most "text-to-SQL" demos are single-shot prompts that hallucinate columns, run unvalidated queries, and are never evaluated. **DataChat** treats the problem as it actually is in production: a multi-step agent that plans, grounds itself in a semantic layer, generates SQL, **guardrails and verifies it before execution**, repairs its own mistakes, and explains the result - with tracing and an evaluation harness proving it works.

Key features

- **Real agent, not a prompt** - LangGraph graph: plan -> retrieve -> generate SQL -> guardrail -> (human approve) -> execute -> verify -> self-repair -> explain -> visualize.
- **Safe by construction** - generated SQL passes an AST guardrail chain *and* runs on a read-only, least-privilege DB role with timeouts. No write path exists.
- **Grounded (RAG-to-SQL)** - a semantic layer (schema docs, synonyms, few-shot examples) is retrieved via pgvector so the model doesn't invent columns.
- **Human-in-the-loop** - approve or edit the SQL before it runs; durable state means the pause survives reloads and cold starts.
- **Swappable providers** - the deployed demo runs on **Groq**; a self-hosted **Ollama** model (behind a token-guarded tunnel) is the same port behind the same router and is one env flag away. A public URL has to answer when my PC is off, so the GPU is a demonstrable capability rather than a dependency. Circuit breaker between them.
- **Instant on repeats** - a normalised whole-answer cache replays a prior answer for the same question, skipping the whole LLM chain, with zero false-positive risk (exact-match, never fuzzy). Measured **~1.0-1.6 s -> ~80 ms** (see Evaluation for conditions).
- **Downloadable reports & data** - every answer can be exported as a Markdown report (question, summary, SQL, table, and links to the source datasets) or a CSV of the result set.
- **Honest out-of-scope answers** - when the governed data has no answer, an optional, injection-hardened web fallback returns a **structured table with a citation on every row**, plus a summary and a downloadable report. Web data is a separate type end to end - its own SSE event, its own report layout, `source_url` in the CSV - so a scrape can never be rendered as verified data, and it never re-enters the SQL path.
- **Evaluated** - a 26-case golden set (21 answerable + 5 that *should* be refused) scored by result-set equality, BIRD-style. A deterministic pipeline gate runs on every PR; the model-quality number is measured separately against a committed baseline. See Evaluation.
- **Observable** - every run and model call traced in MLflow, with a versioned prompt registry.
- **Streaming UI** - watch the agent think; results render as a chart from a backend-emitted Vega-Lite spec.
- **\$0/month** - every component runs on a permanent free tier or your own hardware.

Architecture

A **modular monolith with microservice-grade boundaries** - clean architecture + dependency inversion keep vendors, the DB, and even LangGraph as swappable details. It's microservice-*ready* without paying the ops cost prematurely.

[Diagram] A Mermaid diagram appears here in the source Markdown; it renders on GitHub. See the file named at the start of this chapter.

The differentiators - the **agentic LangGraph core**, the **guardrail + read-only-role defence in depth**, the **provider circuit breaker**, and the **eval harness** - are documented in depth in **Design.md** ([./Design.md](#)) and **TechSpec.md** ([./TechSpec.md](#)). Named patterns (Strategy, Adapter, Decorator, Chain of Responsibility, Circuit Breaker, Template Method, Repository, Builder...) are mapped to SOLID in Design §4.

Tech stack

Layer	Choice	Why
Agent	LangGraph 1.2	Durable state, checkpoints, HITL interrupts, subgraphs
Backend	FastAPI + Pydantic v2 + SQLAlchemy 2 (async)	Typed, async, boundary validation as a control
LLMs	Groq (deployed) + Gemini fallback; self-hosted Ollama swappable behind the same port	Provider abstraction + circuit breaker; a public URL must answer when my PC is off
Data / vectors	Postgres + pgvector (Neon)	One system for relational data <i>and</i> embeddings
Cache / limits	Redis (Upstash)	Result cache, rate limiting, breaker state
LLMOps	MLflow 3.14	Tracing, prompt registry, evaluation, CI gate
Frontend	Next.js 16 + React 19 + Tailwind	Streaming chat; thin Vega-Lite renderer
CI/CD	GitHub Actions	Lint, type, test, security, eval on every PR
Cost	Free tiers only	Runs entirely on \$0 - by design, not by luck

Getting started (local, no API keys needed)

Local dev runs against mocks + seed data, so you can try it before creating any accounts.

Prerequisites: Docker + Docker Compose, and (for non-container dev) uv and pnpm.

```
git clone https://github.com/<you>/datachat.git
cd datachat
cp .env.example .env          # defaults use USE MOCKS=true - no real keys required
docker compose up --build     # postgres+pgvector, redis, backend, frontend, mlflow
```

```
# seed the sample open-data slice + few-shot examples
docker compose exec backend python -m ingestion.run --dataset seed
```

- App: <http://localhost:3000> - API docs: <http://localhost:8000/docs> - MLflow: <http://localhost:5000>

- Run the checks: `make test` - `make lint` - `make sec` - `make eval`

To use real models later, set `USE MOCKS=false` and add keys - see **GOLIVE.md** ([./GOLIVE.md](#)) (generated at the end of the build).

Live deployment

Component	Host	URL
Frontend	Vercel (Hobby)	< https://data-chat-seven.vercel.app >

Component	Host	URL
Backend	Render (free)	<https://datachat-api-wmpd.onrender.com>
Database	Neon (Postgres 16 + pgvector)	-
Cache / rate limit	Upstash Redis	-
Model	Groq llama-3.3-70b-versatile	-

Total cost: \$0/month. Every component is a permanent free tier.

- **Provider choice.** Groq serves the demo rather than the self-hosted Ollama, on purpose: a public URL has to answer when my PC is off, and a quick Cloudflare tunnel changes hostname on every restart. Ollama is the same port behind the same router and is one env flag away - see GOLIVE.md (./GOLIVE.md).
- **Cold-start caveat.** The backend sleeps after ~15 min idle and Neon scales to zero, so the first request after idle takes ~50s and the UI shows a "waking" state. A keep-warm ping every 12 minutes mitigates it. Durable LangGraph checkpoints mean an in-flight human-approval step survives the sleep. This is a documented trade-off of a genuinely \$0 deploy, not a defect.

What I learned / engineering highlights

- **Resilience over free APIs:** a per-provider circuit breaker with Gemini->Groq fallback turns flaky free tiers into a dependable service.
- **Safety as architecture:** an AST guardrail chain plus a read-only least-privilege role means no unsafe SQL can execute even if a layer is bypassed.
- **Evaluation you can trust:** execution-accuracy scoring (result-set equality) gates regressions in CI.
- **Durable agent state:** LangGraph checkpoints let a human-in-the-loop pause survive reloads and cold starts.

What data it answers from

Three datasets ship, selectable at ingestion (`--dataset seed|wdi|owid`). The live demo runs seed.

Dataset	Source	Contents
seed	bundled fixture	The union of the two below - keyless local dev and the deployed demo
wdi	World Bank API (live fetch)	15 countries x GDP per capita, population, life expectancy x 2022
owid	Our World in Data	Same 15 countries x CO ₂ total / per capita / global share x 2021-2022

Four tables: countries, wdi_indicators, wdi_values, owid_co2.

The slice is small on purpose. Neon's free tier is 0.5 GB, and the point of the project is retrieval quality and evaluation, not corpus size. It is real data from real sources - not synthetic - and everything the agent claims is traceable to a row you can download as CSV.

Scaling it is a config change, not a rewrite. The World Bank API exposes on the order of a thousand indicators for ~200 countries; widening means adding entries to `_INDICATORS` and a matching semantic definition. Any new source is a `DatasetConnector` - name plus `async fetch()` -> `RawDataset` - alongside `world_bank.py` and `owid.py`, with no change to the agent, the guardrail or the API.

Where the RAG is

Retrieval does **not** fetch answers. It fetches the *schema* the model needs to write correct SQL, which is what keeps the LLM from inventing column names.

```
question --embed--> pgvector cosine top-k --> only the relevant slice of:
    - table + column descriptions
    - units, synonyms
    - few-shot Q->SQL examples
--> SQL-generation prompt
```

- **Indexed:** each semantic table, each column, and each few-shot example, embedded at ingestion time (`ingestion/steps.py`)
- **Retrieved:** cosine top-k over $\Leftrightarrow$ with an HNSW index (`catalog/pgvector.py`)
- **Used by:** the retrieve node, before plan and generate_sql

Why this reduces LLM usage rather than adding to it: the model never receives the whole schema, only the top-k slice, so prompts stay small and the surface for hallucinated columns shrinks. On top of that, an exact-match answer cache replays repeat questions in ~80 ms with **zero** model calls, and embeddings are computed once at ingestion - a query costs exactly one embedding call.

The curated part matters: indicator *names*, *units* and *descriptions* are written by hand in `ingestion/definitions.py` rather than fetched, so an upstream label change cannot alter the grounding surface. That is a supply-chain path straight into the prompt, closed deliberately.

Evaluation

What produced these numbers

Golden set: **26 cases - 21 answerable + 5 that should be refused** (out-of-scope country, indicators we don't carry, an ambiguous question, a year beyond the data). No golden question duplicates a few-shot example; a test enforces that, because a duplicated question measures copying, not reasoning.

Measured at temperature=0 against the seed dataset. Reproduce with `make eval-real`. Baselines are committed per provider in `backend/eval_baseline.json` (`backend/eval_baseline.json`) - execution accuracy is as much a property of the model as of the pipeline, so one blended number across providers would mean nothing.

Metric	Groq llama-3.3-70b-vrsatile (<i>deployed</i>)	Ollama qwen2.5:7b-instruct (<i>self-hosted</i>)	n
Execution accuracy	0.810	0.667	21
Refusal accuracy	1.00	0.80 ?	5
SQL valid rate	0.952	0.952	21
Explanation faithfulness	0.905	0.857	21

Refusal accuracy was 0.80 and is now 1.00. One in five out-of-scope questions was being answered rather than declined, which on a public demo is a hallucination risk - so it is written up rather than quietly improved. The cause was not a careless model: *"what will global CO₂ be in 2030?"* produced `SUM(co2)/SUM(pop) WHERE year = 2030`, and an aggregate over zero rows returns **one row containing NULL**, so the pipeline saw `row_count = 1` and concluded it had an answer. Fixed by deciding scope deterministically *before* SQL generation (a named country or year outside the loaded slice is a fact, not a judgement) and by treating an all-NULL row as no data. Execution accuracy was unchanged at 0.8095, which is the check that matters - the gate added no false refusals.

_{? The Ollama column was measured **before** the scope gate. Its one refusal miss was *"Show me the best countries"*, an ambiguity case the gate does not catch, so that number is not assumed to have improved and has not been re-measured.}

_{The left column is what the live demo runs. The right is the same set against a self-hosted 7B, kept so the Ollama path stays gated too - the 14-point gap is the cost of running on your own GPU, measured rather than guessed. An earlier README showed 0.80 for a 5-case set, one of whose questions was a verbatim few-shot example; both problems are fixed and the set is now 26 cases with a test that fails the build if leakage returns. Known scoring artifact: result-set equality is strict, so an answer returning country *names* where the gold used ISO codes counts as a miss despite being substantively right - that is 1 of the 4 Groq misses.}

How it is gated

Two tiers, because the published number and the per-PR gate cannot come from the same run - one needs a real model, the other has to be free and deterministic.

Tier	Command	Runs in CI	Gates
Pipeline	<code>make eval</code>	yes, every PR	Graph, retrieval, guardrail, execution and the scorers, with a scripted LLM holding model quality constant. Must score exactly 1.00.
Quality	<code>make eval-real</code>	no (needs a GPU or keys)	Real-model execution accuracy against backend/ <code>eval_baseline.json</code> , tolerance 0.05

Tolerance rationale: one answerable case is worth $1/21 = 0.048$, so 0.05 absorbs a single case flipping and blocks two. Baselines are measured and committed, never hand-written.

Answer-cache latency

End-to-end through the containerised stack (`docker compose up`), real model, five distinct questions asked cold then repeated:

Path	Latency	SSE frames
Cold (full agent chain)	987-1569 ms	15-16
Cached (exact-match replay)	77-87 ms	6-7

_{Conditions: local Ollama with `qwen2.5:7b-instruct` already resident in VRAM, Postgres and Redis on the same host. A cold model load or a free-tier cold start adds tens of seconds to the first number and does not affect the second - so treat this as the steady-state ratio (~15-20x), not a cold-boot claim.}

Known limitations & failure modes

Written from measured runs, not from imagination. Being precise about failure is more useful than another adjective.

The published accuracy understates the system, on purpose

Of the 4 answerable cases Groq misses (17/21 = 0.810), **3 are the same scoring artifact**: the model joins to countries and returns the country *name* where the gold SQL selected the ISO code.

```
gold      SELECT country_iso3, co2_per_capita ...    -> ["QAT", 37.6]
predicted SELECT c.name, o.co2_per_capita ...       -> ["Qatar", 37.6]
```

Result-set equality is exact, so that scores as wrong despite being right - and arguably more readable. A lenient scorer would report roughly **0.95**. The strict number is published anyway, because loosening a metric to flatter yourself is how an eval stops being worth running. The remaining miss is real: *"How many countries are in each income group?"* produced **no SQL at all**, which is also the single point of `sql_valid_rate = 0.952`.

Where the agent mis-routes

- **Ambiguity is handled by the model, not by a rule.** The scope gate settles countries and years deterministically, but *"show me the best countries"* has no named entity to check - it depends on the clarify prompt firing. That is the one refusal case the 7B model got wrong.
- **Indicators are not scope-checked.** The vocabulary is open ("literacy rate", "unemployment"), so a keyword list would refuse phrasings that do work. Those questions still generate SQL, match nothing, and are caught downstream - correct, but it costs a model call the country/year path avoids.
- **The scope gazetteer is not exhaustive.** ~124 country names. A miss is not a wrong answer, only a missed early exit.

Where retrieval is weaker than it looks

Retrieval is over **4 tables** - it is not a hard retrieval problem, and NDCG-style numbers would be meaningless at this size. The harder property is that the semantic layer is **curated by hand**: indicator names, units and descriptions are written in `ingestion/definitions.py` rather than fetched, so an upstream label change cannot alter the grounding surface. That is a deliberate supply-chain decision, but it means widening the corpus is manual work, not an automatic import.

Dev and production also embed differently - a deterministic hash embedder locally, Gemini in production. Changing embedder without re-ingesting silently degrades retrieval, because the stored vectors and the query vector stop sharing a space.

What the eval does not cover

The 26 cases measure single-turn NL->SQL over one seeded slice. **Not covered**: multi-turn conversation, the human-approval path, chart *correctness* (only that a spec is produced), latency, cost, the web-fallback path, or any dataset other than seed. The CI tier holds model quality constant and measures the pipeline; it says nothing about answer quality.

The 7B vs 70B gap is real and measured

	Groq 70B (<i>deployed</i>)	Ollama qwen2.5 7B (<i>self-hosted</i>)
Execution accuracy	0.810	0.667

	Groq 70B (<i>deployed</i>)	Ollama qwen2.5 7B (<i>self-hosted</i>)
Faithfulness	0.905	0.857

A 14-point drop is the price of running on your own GPU on this task. The Ollama refusal figure predates the scope gate and has not been re-measured, so it is not claimed as improved.

Operational

- **~50s cold start** after idle (free tier sleeps, Neon scales to zero).
- **Answer cache is exact-match**, deliberately: "top 5" and "top 10" are nearly identical as strings but need different answers, so fuzzy matching would return a confidently wrong result.
- **Web fallback is off in production**. It works, but it is capped by search *snippets* - typically one or two sparse columns - and enabling it would contradict the published refusal number until the golden set separates "must refuse" from "must escalate".
- **Three cryptography advisories are accepted, not fixed**: mlflow pins cryptography<47. They are ignored by ID in CI with the constraint documented, so they resurface when mlflow relaxes it.

Testing, observability & security

- **Testing** - unit + integration (>=80% on core), agent-eval golden set, Playwright smoke; full matrix in CI. See ImplementationPlan.md (./ImplementationPlan.md).
- **Observability (MLflow)** - 100% of runs traced (spans, tokens, latency, provider, prompt version) + a versioned prompt registry. See TechSpec §8 (./TechSpec.md).
- **Security (OWASP)** - mapped and **tested** against the **LLM Top 10 (2025)** and **Agentic Top 10 (2026)**: prompt-injection corpus, least-privilege tools, read-only execution, rate limits, bounded loops. See ImplementationPlan §11 (./ImplementationPlan.md).

Documentation

Start here: FLOW.md (./FLOW.md) - how a question becomes an answer, end to end, with the trust boundaries marked. - **LEARN.md** (./LEARN.md) - what changed and why, with the alternatives that were rejected.

PRD (./PRD.md) - TechSpec (./TechSpec.md) - AppFlow (./AppFlow.md) - Design (./Design.md) - Schema (./Schema.md) - ImplementationPlan (./ImplementationPlan.md) - Tracker (./Tracker.md) - Rules (./Rules.md) - Security (OWASP matrix) (./SECURITY.md) - Go-live checklist (./GOLIVE.md) - Deploy walkthrough (./DEPLOY.md)

License & data attribution

Code: MIT - see LICENSE (./LICENSE).

Data: not covered by the MIT licence. The bundled seed slice redistributes roughly 90 values from the **World Bank** (World Development Indicators) and **Our World in Data** (CO₂ emissions), both used with attribution under their own terms. Details, links and an accuracy caveat in DATA-SOURCES.md (./DATA-SOURCES.md).

2. Architecture and Flow

How a question becomes an answer: the request lifecycle, the agent graph node by node, and the trust boundaries between governed data, model output and the web. *Source: FLOW.md*

A single-file walkthrough of the system, from an HTTP request to a rendered chart, with the trust boundaries marked. Read this before the code; every section names the files it describes.

Contents

- The 10-second version
- Request lifecycle
- The agent graph
- Node by node
- Trust boundaries
- The SSE event contract
- Where data comes from
- How the LLM call is made
- Caching and reports
- How it is evaluated
- Failure modes
- Deployment topology

1. The 10-second version

A question goes to a **LangGraph agent**. The agent retrieves a *semantic layer* (schema docs + few-shot examples, via pgvector) so the model writes SQL against real columns instead of invented ones. The SQL is parsed to an **AST and guardrailed** before it runs, then executed by a **read-only Postgres role**. The result is verified, explained, charted, and streamed to the browser over **SSE**.

Two things make it defensible rather than a demo: nothing unsafe can execute even if the model is fully compromised, and the whole thing is measured by a golden set with a gate that is calibrated to fail.

2. Request lifecycle

[Diagram] A Mermaid diagram appears here in the source Markdown; it renders on GitHub. See the file named at the start of this chapter.

Files. `interface/api/middleware.py` (backend/app/interface/api/middleware.py) -
`interface/api/rate_limit.py` (backend/app/interface/api/rate_limit.py) - `interface/api/routers/chat.py`
(backend/app/interface/api/routers/chat.py) - `interface/api/sse.py` (backend/app/interface/api/sse.py) -
`application/services/query_service.py` (backend/app/application/services/query_service.py)

Three things worth noticing:

- **The route is** `/api/v1/chat`. Streaming responses use `StreamingResponse(..., media_type="text/event-stream")`.

- Idempotency-Key **is honoured**. A retried POST with the same key does not re-run the agent; it replays a terminal done pointing at the original run.
- **Every error inside the stream becomes a safe error event**. A stack trace never reaches the client - see `_SAFE_MESSAGES` in `query_service.py`.

3. The agent graph

Built by GraphBuilder in `application/agent/graph.py` (`backend/app/application/agent/graph.py`). Solid lines are unconditional edges; diamonds are conditional routers.

[Diagram] A Mermaid diagram appears here in the source Markdown; it renders on GitHub. See the file named at the start of this chapter.

Safety lives in the routing, not in a prompt:

- Nothing reaches `execute` without passing guardrail.
- `hitl_approve` is a **durable server-side interrupt** - the run pauses in the checkpointer and survives a reload or a cold start.
- The repair loop is hard-capped by `max_repair_attempts` (default 2), so it cannot run away. This is OWASP Agentic **ASI08**.
- A user *edit* of the SQL re-enters guardrail, never `execute` directly. Edited SQL is still untrusted input.

State is a TypedDict (`total=False`) so each node returns a partial update that LangGraph merges - see `application/agent/state.py` (`backend/app/application/agent/state.py`).

4. Node by node

Node	Does	Can it interrupt?
understand	Cheap ambiguity check. If the question is vague, <code>interrupt()</code> asks the user to choose.	yes (clarify)
retrieve	Embeds the question once, cosine top-k over semantic tables + few-shot examples via pgvector.	no
plan	Turns the question + retrieved context into ordered steps and target tables.	no
generate_sql	Writes candidate SQL grounded only in the retrieved schema.	no
guardrail	Parses to AST and runs the rule chain. Nothing proceeds on failure.	no
hitl_approve	Durable interrupt: approve / edit / reject the SQL.	yes (approve)
execute	Runs the SQL as the read-only role, with a row cap and statement timeout.	no
verify	Sanity-checks the result against the question.	no
repair	Feeds the failure back for one more attempt, within budget.	no
explain	Prose grounded in the returned rows.	no

Node	Does	Can it interrupt?
visualize	Emits a Vega-Lite spec; the frontend is a thin renderer.	no
web_fallback	Only on an empty result, only when enabled. Search -> attributed table + summary.	no
respond	The single terminal node.	no

Every node inherits `BaseNode`, whose Template Method wraps `_run` in `tracer.span("node.<name>")`.

Tracing coverage is structural - you cannot add an untraced node. See `agent/base_node.py`

(`backend/app/application/agent/base_node.py`).

5. Trust boundaries

This is the part to understand. Three classes of data, never interchangeable.

[Diagram] A Mermaid diagram appears here in the source Markdown; it renders on GitHub. See the file named at the start of this chapter.

Governed (green). Rows produced by SQL that passed the guardrail and ran as `datachat_exec`, a `LOGIN` role with `default_transaction_read_only = on`, a statement timeout, and `SELECT` only on the analytics schema. Defence in depth: even if the guardrail were bypassed entirely, the role cannot write.

Model output (yellow). `candidate_sql` is untrusted until the AST chain clears it: `ReadOnlyRule`, `TableAllowlistRule`, `NoSystemCatalogRule`, `MandatoryLimitRule` (see `infrastructure/sql/validator.py` (`backend/app/infrastructure/sql/validator.py`)). Parsing to an AST - not regex matching - is what makes this hold up.

Web (red). Search snippets and anything derived from them. Kept apart by type, not by a flag:

- `WebTable` is a **distinct domain type**, not an `ExecutionResult`.
- A distinct `web_table` SSE event, so a client cannot render web rows through the governed path by accident.
- A distinct report document with a provenance banner and a per-row `Source` column; the CSV export carries `source_url` per row.
- Web content **never re-enters the SQL path** - enforced by graph topology: `web_fallback` has exactly one outgoing edge, to `respond`.

Why a separate type rather than `ExecutionResult` plus `is_web=True`: one missed flag check silently launders a scrape as verified data, and the whole credibility of the project rests on that distinction holding.

Injection defence is layered, and the prompt is the weakest layer. For the web-table extraction the prompt asks for per-row citations, but `parse_web_table` *enforces* them: any row that is misshapen, all-null, or cites a source outside the ones actually shown to the model is dropped. Prompt instructions are a request; the parser is the control.

6. The SSE event contract

Frames are event: `<type>\n`data: `<json>\n\n`, formatted in `interface/api/sse.py`

(`backend/app/interface/api/sse.py`) from the tagged dataclasses in `agent/events.py`

(`backend/app/application/agent/events.py`).

Event	Payload	When
status	{stage}	Every node transition
plan	{steps, target_tables}	After plan
sql	{sql}	After generate_sql
awaiting_approval	{run_id, kind, sql?, options?}	HITL interrupt (approve or clarify)
rows	{columns, rows, row_count, truncated}	Governed result set
web_table	{columns, rows[{values, source}], row_count, caveat}	Web-sourced table
explanation_delta	{text}	Prose summary
chart_spec	{spec}	Vega-Lite spec
web_sources	{sources[{title, url}]}	Citations for a web answer
error	{code, message}	Safe message, never a trace
done	{run_id, trace_id}	Terminal

A real cold run observed end to end:

```
status status status plan status sql status status rows
status status explanation_delta status status done      (15 frames)
```

A cache replay of the same question is 6-7 frames and ~80 ms.

7. Where data comes from (ingestion)

Offline job, not the request path. `ingestion/` (`backend/ingestion`) is a Chain of Responsibility over `IngestionContext`.

[Diagram] A Mermaid diagram appears here in the source Markdown; it renders on GitHub. See the file named at the start of this chapter.

Two decisions worth defending:

- **Indicator metadata is curated, not fetched.** Values come from the World Bank API; the names, units and descriptions are written by hand in `ingestion/definitions.py` (`backend/ingestion/definitions.py`). The grounding surface cannot be changed by an upstream label edit - that would be a supply-chain path straight into the prompt.
- **Checksum-gated and idempotent.** Re-running ingestion with unchanged data is a no-op (seed: `skipped (unchanged)`), so boot-time seeding is safe.

Current slice (deliberately small - Neon free tier is 0.5 GB): **15 countries**, 3 WDI indicators for 2022, OWID CO₂ for 2021-2022. 4 semantic tables, 4 few-shot examples.

8. How the LLM call is made

[Diagram] A Mermaid diagram appears here in the source Markdown; it renders on GitHub. See the file named at the start of this chapter.

`TaskKind` lets the router pick a provider by strength: `sql_gen` - `repair` - `explain` - `verify` - `clarify` - `classify` - `web_answer` - `web_table`.

Every request-path prompt is **versioned** (sql_generation@v1, explanation@v1, clarify@v1, web_answer@v1, web_table@v1, faithfulness_judge@v1) and the versions used are written into state.prompt_versions and the trace - so any run is reproducible back to the exact prompt text.

temperature defaults to 0.0, which is why the eval's regression tolerance is sized by golden-set granularity rather than sampling noise.

9. Caching and reports

Answer cache - services/answer_cache.py (backend/app/application/services/answer_cache.py). Keyed on sha256(normalised question) where normalising is case/whitespace folding and trailing-punctuation strip, and *nothing more*. Deliberately **exact-match, never fuzzy**: for analytics, "top 5" vs "top 10" and "2022" vs "2021" are nearly identical as text but need different answers, so a fuzzy hit returns a confidently wrong result - strictly worse than a miss.

Measured: **987-1569 ms cold -> 77-87 ms cached** (warm local model).

Reports - GET /api/v1/runs/{run_id}/report.md and /data.csv. Two different documents depending on provenance:

	Governed answer	Web answer
Sections	Summary, SQL , Data, Sources	Warning banner, Summary, Data + Source column, Caveat, numbered links
CSV	columns + rows	columns + rows + source_url
Cached under cache:answer:*?	yes	no - a stale scrape must not replay as fresh
Cached under report:{run_id}?	yes	yes

10. How it is evaluated

Two tiers, because the published number and the per-PR gate cannot come from the same run - one needs a real model, the other must be free and deterministic.

[**Diagram**] A Mermaid diagram appears here in the source Markdown; it renders on GitHub. See the file named at the start of this chapter.

Golden set: **26 cases - 21 answerable + 5 refusal** (services/golden_set.py (backend/app/application/services/golden_set.py)). Refusals are scored separately as refusal_accuracy, because result-set equality against a null gold is meaningless and blending them would make a perfect refusal indistinguishable from a wrong answer.

Measured on qwen2.5:7b-instruct, temperature=0:

Metric	Score	n
Execution accuracy	0.667	21
Refusal accuracy	0.80	5
SQL valid rate	0.952	21
Faithfulness	0.857	21

Structural guards in `tests/unit/test_golden_set.py` fail the build if a golden question duplicates a few-shot example (that measures copying, not reasoning), or if gold SQL could not itself clear the guardrail.

`make eval` (free, CI) - `make eval-real` (needs a GPU or keys).

11. Failure modes and what the user sees

Failure	Behaviour	Surfaced as
Ambiguous question	understand interrupts and offers options	<code>awaiting_approval {kind: "clarify"}</code>
Model writes unsafe/invalid SQL	Guardrail blocks; repair retries within budget	status then a safe error if exhausted
Query times out	Statement timeout on the role	"That query took too long to run - try narrowing it."
Valid query, no rows	Web fallback if enabled, else an honest empty answer	<code>web_table</code> + <code>web_sources</code> , or empty rows
All LLM providers down	Circuit breaker exhausts the chain	"The service is busy right now"
Tracking server down	Tracing degrades to a no-op	nothing - the answer still lands
Free-tier cold start	<code>/ready</code> resumes the DB, keep-warm cron mitigates	a "waking" state in the UI

Known limits, stated plainly. Result-set equality is strict, so an answer returning country *names* where the gold used ISO codes scores as a miss despite being right. Web-fallback quality is capped by search-snippet fidelity - typically one or two sparse columns. And an out-of-corpus question currently burns the full repair budget on SQL before falling through to the web.

12. Deployment topology

[Diagram] A Mermaid diagram appears here in the source Markdown; it renders on GitHub. See the file named at the start of this chapter.

Local `docker compose` mirrors this: `postgres+pgvector`, `redis`, `backend`, `frontend`, `mlflow` - with `USE MOCKS=true` so it runs with no accounts and no keys.

Cold-start caveat, by design not by accident: the free backend sleeps after ~15 min idle and Neon scales to zero, so the first request after idle takes tens of seconds. A keep-warm cron mitigates it; durable checkpoints mean an interrupted HITL run survives the sleep.

Where to look next

Question	File
Why was it built this way?	<code>Design.md</code> (<code>./Design.md</code>), <code>LEARN.md</code> (<code>./LEARN.md</code>)
What are the contracts?	<code>TechSpec.md</code> (<code>./TechSpec.md</code>), <code>Schema.md</code> (<code>./Schema.md</code>)
Sequence diagrams per feature	<code>AppFlow.md</code> (<code>./AppFlow.md</code>)
Threat model + tests	<code>SECURITY.md</code> (<code>./SECURITY.md</code>)
Deploying it	<code>GOLIVE.md</code> (<code>./GOLIVE.md</code>), <code>DEPLOY.md</code> (<code>./DEPLOY.md</code>)

3. Sequence Diagrams

The same flows as message sequences, including human-in-the-loop approval and the out-of-scope path. *Source: AppFlow.md*

Application Flow - end-to-end journeys and lifecycles as sequence diagrams. Reads alongside TechSpec (./TechSpec.md) (§4 API, §5 agent) and Design (./Design.md) (classes).

1. Journeys at a glance

- **Ask** -> **grounded answer** (happy path, streaming).
- **HITL** - **approve/edit SQL** before execution.
- **HITL** - **clarify** an ambiguous question.
- **RAG retrieval** (schema + few-shot grounding).
- **Resilience** - retry -> circuit breaker -> provider fallback.
- **Self-repair loop** on a bad query.
- **Cold-start / keep-warm** (free-tier reality).
- **Ingestion** (offline).
- **Evaluation** (CI gate).

Legend: **FE** Next.js - **BFF** API gateway - **ORCH** LangGraph orchestrator - **SEM** semantic layer - **LLM** provider gateway - **GUARD** SQL guardrail+executor - **PG** Postgres - **CP** checkpointer - **MLF** MLflow.

2. Happy path - ask -> streamed answer

[**Diagram**] A Mermaid diagram appears here in the source Markdown; it renders on GitHub. See the file named at the start of this chapter.

2a. Out-of-scope - the web fallback

Reached only when a **valid** query returned **no rows** and `DATACHAT_WEB_SEARCH_ENABLED=true`. The governed data has no answer, so rather than a dead end the agent searches the web and returns an *attributed* table.

[**Diagram**] A Mermaid diagram appears here in the source Markdown; it renders on GitHub. See the file named at the start of this chapter.

Note the event is `web_table`, **not** `rows` - a client must not render web scrapings through the governed-results path. `web_fallback` has exactly one outgoing edge, to respond, so web content structurally cannot re-enter SQL generation. See FLOW.md §5 (./FLOW.md#5-trust-boundaries).

3. HITL - approve / edit SQL

Interrupt happens **after** the guardrail passes but **before** execution, so a human sees exactly what will run.

[**Diagram**] A Mermaid diagram appears here in the source Markdown; it renders on GitHub. See the file named at the start of this chapter.

Why this matters: the pause is server-side and durable - a user can close the tab, the host can sleep, and the run resumes from the checkpoint. The approval cannot be bypassed by the client (mitigates ASI09 *Human-Agent Trust Exploitation*).

4. HITL - clarify ambiguous question

[**Diagram**] A Mermaid diagram appears here in the source Markdown; it renders on GitHub. See the file named at the start of this chapter.

5. RAG retrieval (grounding)

[**Diagram**] A Mermaid diagram appears here in the source Markdown; it renders on GitHub. See the file named at the start of this chapter.

Only the retrieved subset enters the SQL-gen prompt - small context, narrower surface, fewer hallucinations.

6. Resilience - retry -> circuit breaker -> fallback

[**Diagram**] A Mermaid diagram appears here in the source Markdown; it renders on GitHub. See the file named at the start of this chapter.

If **all** providers are exhausted, the user gets `event: error {code:"providers_unavailable", message:"Busy right now - try again shortly."}` - never a stack trace.

7. Self-repair loop

[**Diagram**] A Mermaid diagram appears here in the source Markdown; it renders on GitHub. See the file named at the start of this chapter.

8. Cold-start / keep-warm

[**Diagram**] A Mermaid diagram appears here in the source Markdown; it renders on GitHub. See the file named at the start of this chapter.

9. Ingestion (offline)

[**Diagram**] A Mermaid diagram appears here in the source Markdown; it renders on GitHub. See the file named at the start of this chapter.

Chain-of-Responsibility pipeline; each step is a link with a uniform `process(context)` contract (see `Design (./Design.md)`).

10. Evaluation (CI gate)

[**Diagram**] A Mermaid diagram appears here in the source Markdown; it renders on GitHub. See the file named at the start of this chapter.

4. Data Model

Tables, roles and grants. The read-only executor role is the second layer of the safety story, independent of the SQL guardrail. *Source: Schema.md*

Data model for both databases-in-one: the app schema (metadata, conversations, semantic layer, eval) and the analytics schema (the open datasets, read-only). Plus pgvector, Redis keys, migrations, seed data, and retention/PII. Companion to Design (./Design.md) and TechSpec (./TechSpec.md). One Neon Postgres instance, two schemas, strict role separation.

1. Overview & isolation model

- app **schema** - everything the product owns: conversations, runs, audit trail, the semantic layer (+ embeddings), eval cases/runs. Read-write from the app's normal role.
- analytics **schema** - the curated open datasets the agent queries. Reachable **only** through a **read-only role** with a statement timeout (see §5). This physical separation is the second layer of the "no writes ever" defence (the guardrail chain is the first).
- **LangGraph checkpoint tables** live in app (auto-managed by langgraph-checkpoint-postgres).

2. ER diagram (app schema)

[**Diagram**] A Mermaid diagram appears here in the source Markdown; it renders on GitHub. See the file named at the start of this chapter.

3. Table details (app schema)

Conventions: uuid PKs (gen_random_uuid()), timestampz in UTC, NOT NULL unless noted, FKs ON DELETE CASCADE for owned children.

Table	Notable columns / constraints	Indexes
conversations	title, user_ref nullable	(updated_at desc)
turns	role CHECK in ('user', 'assistant')	(conversation_id, created_at)
runs	status CHECK, prompt_versions jsonb, trace_id	(conversation_id, created_at), (status)
agent_actions	append-only audit/outbox; never updated	(run_id, created_at)
datasets	name UNIQUE, checksum, version	(name)
semantic_tables	embedding vector(768)	HNSW on embedding
semantic_columns	unit, synonyms jsonb, embedding vector(768)	HNSW on embedding, (semantic_table_id)
few_shot_examples	question, sql, embedding vector(768)	HNSW on embedding
eval_cases	gold_sql	(dataset_id)
eval_runs	metrics + mlflow_run_id	(created_at desc)
eval_case_results	passed, predicted_sql, failure_reason	(eval_run_id)

4. Analytics schema (curated open data, read-only)

Bounded on purpose (Neon free = 0.5 GB). v1 loads a slice of **World Bank WDI + Our World in Data (CO₂)**. Star-ish shape: dimension countries, catalog wdi_indicators, fact wdi_values.

[**Diagram**] A Mermaid diagram appears here in the source Markdown; it renders on GitHub. See the file named at the start of this chapter.

Constraints: wdi_values PK (country_iso3, indicator_code, year); owid_co2 PK (country_iso3, year); FKs to countries. Indexes on (indicator_code, year) and (country_iso3, year) for the common filter/aggregate patterns. All numeric measures nullable (real-world gaps - this is exactly what BIRD-style "dirty real data" tests).

5. Read-only role & execution safety (the security spine)

```
-- Roles: app_rw (migrations + app CRUD on app schema) and analytics_ro (query execution).
CREATE ROLE analytics_ro NOLOGIN;
GRANT USAGE ON SCHEMA analytics TO analytics_ro;
GRANT SELECT ON ALL TABLES IN SCHEMA analytics TO analytics_ro;
ALTER DEFAULT PRIVILEGES IN SCHEMA analytics
    GRANT SELECT ON TABLES TO analytics_ro;
REVOKE ALL ON SCHEMA app FROM analytics_ro;    -- cannot touch app data

-- A dedicated login user for the executor connection pool, in the RO role:
CREATE ROLE datachat_exec LOGIN PASSWORD 'exec_pw' IN ROLE analytics_ro;
ALTER ROLE datachat_exec SET statement_timeout = '5s';
ALTER ROLE datachat_exec SET default_transaction_read_only = on;
ALTER ROLE datachat_exec SET search_path = analytics;
```

The **executor uses its own engine/pool** as datachat_exec (Bulkhead). Even if every guardrail were bypassed, the database itself would reject any write, any cross-schema read, and any query over 5s. Defence in depth (FR-23, LLM06, ASI02/ASI03).

6. pgvector strategy

- **Dimension:** 768 (Gemini text-embedding-004; if the embedding provider changes, dimension is a config + migration change).
- **Index:** HNSW with vector_cosine_ops (pgvector 0.8) - good recall/latency for our small corpus; IVFFlat documented as the low-memory alternative.
- **What's embedded:** table descriptions, column descriptions (+ synonyms/units), and few-shot question text. Retrieval = cosine top-k over each set.
- **Query:** SELECT ... ORDER BY embedding <=> :q LIMIT :k. Because the corpus is tiny, index build/maintenance cost is negligible.

7. Redis (Upstash) key patterns & TTLs

Budget-aware: 500k commands/month. Cache only hot paths; short TTLs.

Key	Value	TTL	Purpose
rl:{scope}:{id}:{window}	counter	= window (e.g. 60s)	Rate limit per IP/session (token/fixed window)
quota:global:{yyyymmdd}	counter	24h	Global daily LLM-call cap (protects free tiers)

Key	Value	TTL	Purpose
cache:answer:{sha256(dataset+question)}	answer JSON	1h	Skip the whole agent run on repeat questions
cache:sql:{sha256(sql)}	result JSON	15m	Skip re-execution of identical SQL
idem:{idempotency_key}	run_id	24h	De-dupe retried POST /chat
breaker:{provider}	state + opened_at	cooldown (e.g. 30s)	Shared circuit-breaker state across restarts

No user content beyond the transient cache lives in Redis; nothing sensitive is persisted there.

8. Migrations (Alembic)

- Async Alembic env; **autogenerate** for the app schema from SQLAlchemy models.
- **Hand-written migrations** for: enabling pgvector (CREATE EXTENSION IF NOT EXISTS vector), creating the analytics schema, and the **roles/grants** in §5 (security-critical -> explicit, reviewed, not autogenerated).
- The analytics **data** is populated by the ingestion job, not migrations (migrations own structure + roles; ingestion owns rows).
- Every migration is reversible; CI runs upgrade head then downgrade -1 on a scratch DB.

9. Seed data (for local dev without any keys)

- A tiny fixture slice (tests/fixtures/analytics_seed.sql): ~30 countries, ~10 WDI indicators, a few years, and the OWID CO₂ subset - enough to answer the golden set locally.
- few_shot_examples and eval_cases seed files so docker compose up yields a working, testable app with **mock** LLM/embeddings (FR-25).

10. Retention & PII

- **No PII in datasets** - public, aggregate open data only (LLM02).
- **User questions** (turns.content) are user-authored text; retained **90 days** by default, then purged by a scheduled job (scripts/purge_old_turns.py). Configurable via DATA_RETENTION_DAYS.
- **IP addresses** are used only for rate limiting and live **only in Redis** with short TTLs - never written to Postgres. If logged for abuse investigation, they are **hashed**.
- **Audit trail** (agent_actions) keeps SQL + decisions (no PII) for debugging/security; same 90-day window.
- Secrets never touch the DB or logs (see Rules (./Rules.md)).

5. Security

OWASP LLM Top 10 and Agentic Top 10, each mitigation mapped to a passing test rather than to a claim. *Source: SECURITY.md*

DataChat is built untrusted-by-default: **every user input and every LLM output is validated at a boundary**, and no single layer is trusted to catch everything. This document maps the OWASP **LLM Top 10 (2025)** and **Agentic Top 10 (2026)** to the concrete mitigation in the code and the test that proves it.

Run the suite: `cd backend && uv run python -m pytest -m security` (plus the DB-backed cases under `pytest -m "integration or eval"` in CI).

Defence-in-depth spine

- **Guardrail (app layer):** sqlglot-AST validator chain - single statement, read-only, table allow-list, no system catalogs, mandatory LIMIT.
- **Read-only DB role (data layer):** `datachat_exec` - SELECT-only on the analytics schema, `statement_timeout`, `default_transaction_read_only`, no access to the app schema.

Either layer alone blocks a write; both together is the guarantee.

OWASP LLM Top 10 (2025)

Risk	Mitigation	Test
LLM01 Prompt Injection	Untrusted input + data cells; hardened system prompt ("schema is data, not instructions"); guardrail before execute; a compromised model output still can't write. Web-search fallback: snippets are fenced in <code><results></code> and the prompt forbids following instructions in them; web content never reaches the SQL path - enforced by graph topology, <code>web_fallback</code> has exactly one outgoing edge (to respond)	<code>test_owasp_llm.py::test_llm01_compromised_model_output_cannot_write</code> , <code>test_sql_injection_corpus.py</code> , <code>test_web_fallback.py::test_web_answer_prompt_wraps_untrusted_content_as_data</code>
LLM01a Injected rows via structured web extraction	Extracting a <i>table</i> from untrusted snippets widens LLM01: a hostile page can try to inject a row that renders in a data table and a downloadable report, with a fabricated citation to look sourced. The prompt requires per-row attribution, but <code>parse_web_table</code> enforces it - rows that are misshapen, all-null, or cite a source index outside the results actually shown to the model are dropped, and columns/rows are capped. Nulls render as an em dash, never "None", so a gap cannot be misread as data. Prompt instructions are a request; the parser is the control	<code>test_web_table.py::test_drops_rows_citing_a_source_that_was_never_shown</code> , <code>::test_drops_rows_whose_width_does_not_match_the_columns</code> , <code>::test_nested_structures_are_not_treated_as_cells</code> , <code>::test_prompt_marks_the_snippets_as_untrusted_and_numbers_them</code>

Risk	Mitigation	Test
LLM09 Misinformation via provenance confusion	Web-derived rows are a distinct domain type (WebTable, not ExecutionResult), a distinct SSE event (web_table, not rows), and a distinct report document with a warning banner and a per-row Source column; the CSV export carries source_url so provenance survives leaving the app. A single shared type with a provenance flag was rejected: one missed flag check silently presents a scrape as guardrailed data	test_web_table.py::test_web_report_is_labelled_and_cites_eververy_row
LLM02 Sensitive Info Disclosure	Public non-PII data only; SecretStr keys; log minimization; safe error messages	test_secret_hygiene.py, test_api.py::test_provider_outage_becomes_safe_error_event
LLM03 Supply Chain	Pinned deps + uv.lock/pnpm-lock; pip-audit/pnpm audit in CI	test_owasp_llm.py::test_llm03_dependencies_are_pinned_with_a_lockfile, CI audit job
LLM04 Data & Model Poisoning	Curated sources; ingestion validates shape + checksum; grounding content is curated, never fetched	test_ingestion_pipeline.py::test_tampered_data_is_rejected_by_checksum
LLM05 Improper Output Handling	Every LLM output validated; SQL guardrailed; chart is validated JSON, not code	test_owasp_llm.py::test_llm05_*, test_charts.py
LLM06 Excessive Agency	Read-only role; fixed least-privilege node/tool set; no dynamic tool loading; HITL before execute	test_owasp_llm.py::test_llm06_tool_set_is_fixed_no_dynamic_loading, test_readonly_role.py
LLM07 System Prompt Leakage	No secrets in system prompts; assume leakable	test_secret_hygiene.py::test_system_prompt_leakage_reveals_nothing_sensitive
LLM08 Vector/Embedding Weakness	Only curated docs embedded; retrieval read-only; poisoned source rows rejected at ingest	test_ingestion_pipeline.py::test_unexpected_table_or_column_is_rejected
LLM09 Misinformation	Grounding + verify node + row-cited explanation + faithfulness scorer	test_eval.py::test_faithfulness_judge_parses_score, golden eval
LLM10 Unbounded Consumption	Rate limits + global quota; capped retries/repair; timeouts; circuit breaker	test_owasp_llm.py::test_llm10_repair_loop_is_bounded, test_api.py::test_rate_limit_returns_429_with_retry_after, test_llm_decorators.py

OWASP Agentic Top 10 (2026)

Risk	Mitigation	Test
ASI01 Agent Goal Hijack	Scoped system role; retrieved content is data; bounded read-only action space	test_owasp_agentic.py::test_asi01_goal_hijack_does_not_change_the_action_space

Risk	Mitigation	Test
ASI02 Tool Misuse & Exploitation	Fixed tool set; read-only executor; SQL argument validation	<code>test_owasp_agentic.py::test_asi02_tool_arguments_are_validated</code> , <code>test_sql_injection_corpus.py</code>
ASI03 Identity & Privilege Abuse	Separate <code>datachat_exec</code> RO role (Bulkhead); no app-schema access	<code>test_readonly_role.py::test_readonly_role_cannot_read_app_schema</code>
ASI04 Agentic Supply Chain	Pinned deps; scanners; verified provider endpoints (constants)	CI supply-chain job, <code>test_llm03_*</code>
ASI05 Unexpected Code Execution	No <code>eval/exec/compile/shell</code> ; SQL only via guardrail + RO role; <code>chart</code> = declarative JSON	<code>test_no_dynamic_execution.py</code>
ASI06 Memory & Context Poisoning	Durable checkpoint integrity; curated few-shots; retrieved data sandboxed as reference	<code>test_owasp_llm.py::test_llm01_*</code> (data-as-instructions refused)
ASI07 Insecure Inter-Agent Comms	Out of scope by design - single process, one StateGraph, no external agents	<code>test_owasp_agentic.py::test_asi07_inter_agent_comms_are_out_of_scope_by_design</code>
ASI08 Cascading Agent Failures	Hard caps on repair/retries; circuit breakers; timeouts	<code>test_owasp_llm.py::test_llm10_repair_loop_is_bounded</code> , <code>test_circuit_breaker.py</code>
ASI09 Human-Agent Trust Exploitation	Server-side, non-bypassable HITL; the exact SQL is shown before it runs	<code>test_owasp_agentic.py::test_asi09_human_approval_is_not_client_bypassable</code>
ASI10 Rogue Agents	Bounded action space (read-only, fixed tools); append-only <code>agent_actions</code> audit of every executed query	<code>test_owasp_agentic.py::test_asi10_every_executed_query_is_audited</code>

Scanners (CI)

- `gitLeaks` - zero secrets in the repo/history (pre-commit + CI).
- `bandit` + `ruff` `flake8-bandit` (S) - SAST, zero high/medium.
- `semgrep` - SAST in CI (Linux).
- `pip-audit` + `pnpm audit` - dependency advisories, zero unresolved.

6. Change Log and Reasoning

What changed, why, and what was rejected. Includes the three false claims found in this repo and how each was fixed. *Source: LEARN.md*

A running record of non-obvious changes: what moved, **why**, what else was considered, and how to defend it out loud. Most recent date first; within a date, in the order the work happened.

For *how the system works* rather than why it changed, see FLOW.md (*./FLOW.md*).

2026-08-13 - Live

<https://data-chat-seven.vercel.app/> -> <https://datachat-api-wmpd.onrender.com>

Groq on the free tier, Neon Postgres + pgvector, Upstash Redis, Vercel frontend, keep-warm every 12 minutes. Verified in the browser: "Which 3 countries had the lowest life expectancy in 2022?" -> Nigeria 53.6, South Africa 62.3, India 67.7, with the executed SQL, a Vega-Lite chart and a grounded explanation.

Four failures, and what each really was

None of them said what it meant, which is the whole lesson.

1. `render.yaml` not found. The file had never been pushed. Ten commits - including the one that created it - had been sitting local since 25 July. The GitHub repo a recruiter would have seen was missing the web fallback, the downloadable reports and the answer cache too.
2. Exited with status 127. Render runs `dockerCommand` through a shell, so an inline `sh -c "a && b && c"` had its quotes passed through literally and the shell went looking for one command named `a && b && c`. Moved to `backend/start.sh` - no quoting ambiguity, and readable locally. Took `exec uvicorn` and set `-e` with it, so the server gets `SIGTERM` directly and a failed migration stops the container rather than serving against a half-migrated DB.
3. error parsing value for field `"cors_origins"`. `pydantic-settings` JSON-decodes collection-typed fields at the *source*, before validators run, so a plain `http://localhost:3000` is a JSON syntax error. Decoding off + a validator that takes plain, comma-separated, or JSON.
4. **UI live, every request blocked.** `CORS_ORIGINS` was `https://data-chat-seven.vercel.app/` - the trailing slash you get by copying a URL out of a dashboard. A browser origin is scheme + host + port and never has a path, and Starlette compares exactly. The failure is one-sided and vicious: the API still answers 200 to curl while the browser blocks it, so it reads as "the frontend is broken" and sends you into the wrong codebase.

why fix 3 and 4 in code rather than in the runbook: both arrive by default - the slash from copy-paste, the JSON from a type annotation nobody sees - and neither symptom points at its cause. A note in *GOLIVE.md* would have been a note about a trap rather than the removal of one.

Earned its keep

The MLflow startup bound shipped two days earlier fired in production on the first successful boot: `prompt_register_timed_out` after 5.0s. Without it, startup blocks ~90s on an unreachable tracking server and Render's health check fails the deploy. It was written for exactly this and never tested against it until now.

The `no_llm_provider_configured` guard also paid off - its *absence* from the logs was how I confirmed the Groq key had landed, rather than discovering it later through identical wrong answers.

Deliberately not deployed

Ollama. A public URL must answer when the author's PC is off, and `cloudflared tunnel --url` issues a new hostname on every restart, so every reboot would mean editing the `Render` env and redeploying. It stays one env flag away for a live demo - which is a better interview moment than a link that happens to work.

2026-08-11 - Making the evaluation honest

The problem, stated plainly

Three things in this repo were false or misleading. All three were public.

1. The CI "regression gate" could not fail. `tests/agent_eval/test_golden_eval.py` asserted:

```
assert 0.0 <= report.execution_accuracy <= 1.0
```

That is a tautology. Worse, the job it guarded ran `MockLLMProvider`, which returns *one hardcoded CO₂ query for every question* regardless of what was asked (`app/infrastructure/llm/mock.py`). Measured: the CI eval scored **execution_accuracy = 0.0** and passed. Meanwhile `EvalReport.regressed()` existed and was never called outside unit tests.

The README said the golden set "gates every change in CI." It gated nothing.

2. Two metric names, one number. `guardrail_pass_rate` and `sql_valid_rate` were computed from the identical expression. The README printed both as 1.00, which reads as two independent confirmations of safety. It was one measurement counted twice.

3. Train/test leakage in the golden set. `golden_set.py` claimed "questions are distinct from the few-shot examples to avoid train/test leakage." One of its five questions - "*What was Germany's CO₂ per capita in 2022?*" - was a **verbatim** few-shot example from `ingestion/definitions.py`. The retriever puts that example straight into the prompt for that question, so the model copies the answer. It passed, and inflated the published score.

What changed

Area	Before	After
CI gate	<code>assert 0.0 <= acc <= 1.0</code>	Scripted-LLM pipeline gate, must score exactly 1.00
Quality gate	none	<code>make eval-real</code> vs committed <code>eval_baseline.json</code> , tolerance 0.05
Golden set	5 cases, 1 leaked	26 cases (21 answerable + 5 refusal), leakage blocked by a test
Refusals	not represented	Scored separately as <code>refusal_accuracy</code>
<code>guardrail_pass_rate</code>	duplicate of <code>sql_valid_rate</code>	deleted
Published accuracy	0.80	0.667 - measured, harder set, no leakage

Decision 1 - two tiers, not one gate

The published quality number and the per-PR gate cannot come from the same run. One needs a real model (GPU or paid keys, minutes per run, mild nondeterminism); the other must be free and deterministic on every pull request. Forcing them together is what produced the fake gate in the first place.

- **Tier 1**, `pytest -m eval - real graph, real pgvector catalog, real read-only executor, but a GoldScriptedLLM that returns the gold SQL for each question. Model quality is pinned at perfect, so the score moves only when the pipeline moves. Must be exactly 1.00. Free, runs in CI.`
- **Tier 2**, `pytest -m eval_real - same set, real provider, compared against a committed baseline. Opt-in via DATACHAT_EVAL_REAL=1.`

Alternatives considered. Gate CI on a real model - needs paid keys and tolerates flakiness, and CI has no GPU. Publish the mock number - meaningless, as demonstrated. Keep one blended gate - that is the status quo that failed.

Defence in two sentences: the CI gate holds model quality constant so it measures the pipeline, and it is calibrated to fail - I verified that by mutating the scripted SQL and watching it drop off 1.00. The model-quality number comes from a separate, opt-in run against a committed baseline, so neither number pretends to be the other.

Decision 2 - refusals scored separately, not blended in

Five cases have no correct answer in the governed data. Result-set equality against a null gold is meaningless, so they are scored on whether the agent *declined*: `refusal_accuracy`, over the refusal cases only. Execution accuracy is computed over answerable cases only.

Why it matters: an agent that confidently answers an unanswerable question is the failure mode that actually hurts a user, and blending refusals into one average makes a perfect refusal and a wrong answer indistinguishable.

Alternative considered: score refusals as `execution_accuracy = 0`. Simpler, but it destroys exactly the signal the negative cases exist to produce.

A refusal is detected as: an error/guardrail dead-end, a clarify interrupt (the run pauses with no execution), or a query that ran and matched nothing.

Decision 3 - regression tolerance of 0.05

A golden set is granular. With 21 answerable cases, one case flipping moves the score by $1/21 = 0.048$. So:

- tolerance < 0.048 -> any single flaky case fails the build. That is noise, not regression.
- tolerance **0.05** -> absorbs exactly one case, blocks two (0.095).
- tolerance > 0.095 -> stops catching real slides.

Sampling variance is not the driver - temperature is 0.0 - so the tolerance is sized by set granularity, not by model jitter. If the set grows, revisit it: at 50 answerable cases one case is 0.02 and the tolerance should tighten.

Decision 4 - deleted `guardrail_pass_rate` rather than inventing a difference

It measured the same thing as `sql_valid_rate`. The options were to fabricate a distinction or to delete one name. Since nothing unsafe can execute by construction (AST guardrail *and* a read-only role), there is no second number to report; a genuinely different metric would need a fault-injection setup that does not exist. Deleted. The `eval_runs.guardrail_pass_rate` DB column is retained and now receives `sql_valid_rate` -

renaming it needs a migration and buys nothing.

The number went down. That is the point.

0.80 -> 0.667. Nothing regressed. The old figure came from five easy cases, one of which the model could copy out of its own prompt. The new figure comes from 26 cases spanning lookups, aggregations, rankings, joins, group-by, two-year time-series and out-of-scope refusals, with leakage blocked by a test.

A number that only moves up is not a measurement.

Unplanned fix found along the way

OpenAICompatibleAdapter sent Authorization: Bearer {key} unconditionally. With no key configured - the default (ollama_api_key = "") - httpx rejects the bare "Bearer " as an illegal header value, and the failure surfaced as transport error: Illegal header value b'Bearer ' rather than a config problem. This blocks a keyless local Ollama, which is the documented go-live path in GO_LIVE.md. The header is now omitted when there is no key; the tunnel still enforces auth in production.

Verification performed

- 228 unit/security tests pass (was 219; 9 added).
- Tier-1 gate passes against live Postgres, and **fails when mutated** - the scripted SQL was deliberately corrupted, the gate reported pipeline regression: [...] and exited 1. Reverted after.
- Real-model measurement run twice: 5-case set reproduced the old 0.80 exactly (confirming the previous figure was honest, just easy), 26-case set produced the new baseline.
- 7 integration tests fail locally on datachat_exec password auth. Confirmed pre-existing by stashing all changes and re-running on a clean tree: **same 7 failures**. Environmental (throwaway Postgres volume), not caused by this work.

2026-08-11 - Container verified end to end (C1)

Both images build clean. The full five-service stack (postgres, redis, backend, frontend, mlflow) comes up from a wiped volume, auto-migrates and auto-seeds on boot, and /ready returns 200.

SSE verified against the real model, not mocks. The container ran with USE MOCKS=false, Ollama enabled, qwen2.5:7b-instruct. POST /api/v1/chat returns content-type: text/event-stream and streams status -> plan -> sql -> rows -> explanation_delta -> done - 15 frames for a cold answer, 6-7 for a cache replay. Answers were correct against the seed slice (top-5 CO₂ per capita: Qatar, Saudi Arabia, Australia, US, Canada; highest-average region: Middle East & North Africa at 28.15).

Live confirmation of a known scoring artifact: the agent returned country *names* where the golden SQL uses ISO codes - substantively right, scored as a miss. That is exactly the case flagged in the notes field of that golden case.

Note the route prefix is /api/v1/chat, not /chat.

Cache figure restored, corrected

The old "~20-40 s -> ~70 ms" had no artifact behind it and was cut. It is now measured and back: **987-1569 ms cold -> 77-87 ms cached** across five distinct questions, roughly 15-20x. The old *cached* number (~70 ms) was accurate; the old *cold* number was not reproducible on a warm local model - it likely came from a cold-start context that includes model load. The README now states the conditions so the ratio can't be mistaken for a cold-boot claim.

Still open

- **7 integration tests fail on this host** (`test_readonly_role.py`, `test_executor.py`), all on `InvalidPasswordError` for `datachat_exec`. Proven pre-existing: stashing every change and re-running on a clean tree reproduces the same 7. Also fails in isolation, so it is not test-ordering, and it survives a wiped volume. The role is created with a real SCRAM verifier - checked `pg_authid`, so this is **not** an empty-password hole - but the verifier does not match what the host-side test connects with. Suspect the `set_config('datachat_exec_pw', ...)` -> `current_setting(...)` handoff in `0003_roles_grants` does not resolve as intended when alembic is driven from the host. CI (Linux, container-run migrations) is the reference environment. Worth its own session; untouched here.
- Backend image is 3.16GB; ~1.13GB is a duplicated venv layer caused by `chown -R` after `COPY`, ~470MB is MLflow's scientific stack, and dev tools (`pytest`, `mypy`, `ruff`, `bandit`) ship to production because `uv sync` lacks `--no-dev`. Matters for a Render free-tier cold start.

2026-08-11 - Structured answers from the web fallback

The problem

"Not in the available datasets" is a dead end. The fallback existed but returned a paragraph, and most of these questions ("which countries have X", "what is Y in Z") are inherently tabular - a paragraph reads badly and cannot be exported.

What changed

`web_fallback` now makes a second versioned LLM call (`web_table@v1`) that extracts a small table from the search snippets, rendered with per-row provenance and downloadable as a report or CSV.

Decision 1 - a separate type, not a flag

`WebTable` is its own domain entity. It is *not* an `ExecutionResult` with `is_web=True`.

Governed rows are verified by a read-only role and an AST guardrail. Web rows are a language model's reading of untrusted snippets. Give them one type and they become substitutable in the UI, the report, the cache and the chart - and one missed flag check silently launders a scrape as verified data. The separation is carried all the way out: a distinct `web_table` SSE event, a distinct report document, a `source_url` column in the CSV.

Alternative considered: one type plus a provenance flag. Fewer types, but the failure mode is invisible and permanent, and the credibility of this whole project rests on that distinction holding.

Defence in two sentences: governed data and web scrapings are different types end to end, so the UI cannot render one through the other's path by accident. The web path has exactly one outgoing edge in the graph - to respond - so web content structurally cannot re-enter SQL generation.

Decision 2 - the parser enforces attribution, not the prompt

Once extraction is structured, the obvious injection payload changes. A hostile page no longer just skews a sentence; it tries to inject a row that renders in a data table and a downloadable report, with a fabricated citation to look sourced.

So `parse_web_table` re-validates everything the model returns and drops any row that is misshapen, all-null, or cites a source index outside the results actually shown to it. Columns and rows are capped. Nulls render as an em dash, never the string "None", so a gap cannot be misread as data.

Prompt instructions are a request. The parser is the control. There is a test for the fabricated-citation case specifically.

Decision 3 - downloadable, but never replayed

Web answers are written to report: {run_id} so the user can download the answer they just got, but deliberately **not** to the question-keyed answer cache. The web moves; serving a month-old scrape as a fresh answer is worse than a cache miss. This preserved the existing "don't cache web answers" decision while fixing the fact that reports were silently unavailable for every web answer.

Verified

End to end in the container against live DuckDuckGo and qwen2.5:7b-instruct. "What is the adult literacy rate in Kenya?" produced a cited table (82%, 2020, attributed to Africa Check), a report with the provenance banner, and a CSV carrying source_url. 240 unit tests pass; the tier-1 eval gate still passes.

Honest limits

Snippet fidelity is the ceiling, and it is low. Typical output is one row and one or two columns, one of which is often just "Year" - the Kenya CO₂ answer was a single cell. This was a deliberate choice against adding a page fetcher, which would give real tables at the cost of a new dependency, SSRF and injection surface, latency, and robots/ToS obligations on a public demo.

Consequence to handle before enabling in production

The feature is off by default (DATACHAT_WEB_SEARCH_ENABLED=false) and the eval graph is built without the fallback, so refusal_accuracy = 0.80 still holds.

The moment it is enabled, three of the five refusal cases become wrong - Kenya's CO₂, India's literacy rate and Germany's unemployment would escalate to the web rather than refuse, which is now the *desired* behaviour. The golden set needs splitting into "must refuse" (genuinely ambiguous) and "must escalate" (out-of-corpus but answerable), with a separate escalation_accuracy. Not done.

Also found

- The repair loop burns the full budget - three SQL generations - on questions that can never be answered, before falling through to the web. Real latency and cost per out-of-corpus question.
- Backend startup blocks on MLflow. With the tracking server down, boot took ~90s instead of seconds: the tracer is best-effort at *runtime* but not at *startup*. That will hurt on a free-tier cold start.